
The Netdot Manual

Network Documentation Tool

Version 0.9

Copyright © 2009 University of Oregon

Table of Contents

The Network Documentation Tool	1
Structure	2
Installation	2
Requirements	2
Configuration	3
Upgrading	3
Installing the package for the first time	3
Operation	5
Device Management	5
IP Address Space	6
DNS	7
DHCP	8
Contact Information	8
Cable Plant	8
Advanced DB operations	9
Reports	9
Exporting Configurations for External Programs	9
Authorization	9

This manual documents the installation, administration and operation of the Netdot application.

The Network Documentation Tool

Netdot [<http://netdot.uoregon.edu>] is an open source tool designed to help network administrators collect, organize and maintain network documentation.

Netdot is actively developed by the Network and Telecommunication Services [<http://ns.uoregon.edu>] group of the University of Oregon [<http://www.uoregon.edu>].

Netdot features include:

- Device discovery via SNMP
- Layer 2 topology discovery and graphing, using multiple sources of information: CDP+LLDP, Spanning Tree Protocol, switch forwarding tables, router point-to-point subnets.
- IPv4 and IPv6 address space management (also referred to as *IPAM*), including hierarchical organization, address block visualization and IP and MAC address location and tracking.
- Cable plant information including: sites, rooms, jacks, closets, inter and intra-building wiring, circuits, etc.
- Contact information for related entities: departments, providers, vendors, etc.

- Netdot can generate configuration files for various other tools, including: Nagios [<http://www.nagios.org>], Sysmon [<http://www.sysmon.org>], RANCID [<http://www.shrubbery.net/rancid>], Cacti [<http://www.cacti.net>]. Starting with version 0.9, Netdot also generates configurations for ISC BIND and ISC DHCPD.

Structure

Netdot consists of several components:

1. The database

Our goal has been to make Netdot database-agnostic as much as possible. In principle, it should be able to use any database supported by Perl DBI. There are, however, some limitations to this, for example, schema migration scripts are db-specific and may not always be available. Currently MySQL is fully supported. There is currently partial support for PostgreSQL.

2. The libraries

The back-end code is a hierarchy of object-oriented Perl classes. It can function as an API as well. One advantage of this model is that presentation, collection and database can be separated among different physical machines.

3. User Interface (UI)

The web user interface is built on a templating system called HTML::Mason.

4. Command Line scripts (CLI)

Certain tasks, like device discovery, can be executed from the command line. Therefore, these tasks can be automated by running them periodically via CRON.

Installation

Requirements

- Perl 5.6.1 or later
- Apache2 with mod_perl2
- MySQL (5.x preferred) or PostgreSQL (limited support)
- Authentication Server. Netdot supports both Radius and LDAP.
 - For Radius, we recommend FreeRadius, available at:
<http://www.freeradius.org>
 - For LDAP, you can try OpenLdap, available at:
<http://www.openldap.org>
- The RRDtool package, including its Perl modules, available at:
<http://oss.oetiker.ch/rrdtool/>
- The GraphViz package, available at:

<http://www.graphviz.org>

- Various Perl modules. To test for missing modules in your system, run:

```
%make tesdeps
```

Tip

If any of these modules are provided by your distribution, we recommend installing those.

- You can install any missing modules automatically using the CPAN, by running:

```
%make installdeps
```

If you want to install modules individually, you can do this instead:

```
%perl -MCPAN -e shell  
>install Module::Blah
```

Tip

If you have problems installing Net::IRR, it is possible that the Net::TCP module cannot be installed. There is an open bug and patch available at:

<http://rt.cpan.org/Public/Bug/Display.html?id=43071>

Configuration

Netdot comes with a configuration file that you need to customize to your needs. You need to:

```
% cp etc/Default.conf etc/Site.conf
```

Then, modify Site.conf to reflect your specific options. The original file contains descriptions of each configuration item.

Netdot will first read Default.conf and then Site.conf

The reason for keeping two files is that when an upgrade is performed, the Default.conf can be re-written (to add new variables, etc.), without overwriting your site-specific configuration.

Upgrading

Netdot's database schema only changes between major versions. For example, if upgrading from 0.8.x to 0.9.x, you will need to run an upgrade script to adapt your current database to the new schema. This can be usually accomplished by running the 'make upgrade' command. However, if upgrading from 0.8.X to 0.8.Y, it should suffice to run 'make install' and restart Apache.

Installing the package for the first time

- Make sure you have created the file etc/Site.conf with your configurations (See above).
- From the top directory in the package, do:

```
%make install [parameters]
```

Possible parameters include:

```
PREFIX=YOUR-PREFIX (default: /usr/local/netdot)
APACHEUSER=USER-YOUR-APACHE-RUNS-AS (default: apache)
APACHEGROUP=GROUP-YOUR-APACHE-RUNS-AS (default: apache)
```

- You will then be ready to initialize the database.

```
%make installdb [parameters]
```

Remember you need to set DB_DBA and DB_DBA_PASSWORD to your database's admin username/password in Site.conf before running this command. Or if you prefer, you can specify the DB_DBA and DB_DBA_PASSWORD values as parameters (DB_DBA and DB_DBA_PASSWORD are used by many functions in Netdot, they will need to be set to the correct value in Site.conf eventually)

```
DB_DBA=DATABASE-ADMIN-ACCOUNT
DB_DBA_PASSWORD=DATABASE-ADMIN-PASSWORD
```

Apache Configuration

Edit the supplied Apache config template for either Radius or LDAP, copy it to your Apache config dir and include it somewhere in your Apache config (e.g.):

```
Include conf/netdot_apache2_<radius|ldap>.conf
```

Alternatively, some Apache environments provide a directory from which files are included automatically when Apache starts. In that case, it's sufficient to copy the file into that directory, for example:

```
# cp /usr/local/netdot/etc/netdot_apache2_radius.conf /etc/apache2/conf.d/
```

Tip

Make sure you use the version of the file that gets copied into your install directory by 'make install'. This file contains relevant path substitutions based on your chosen install prefix

CRON jobs

Netdot comes with a few scripts that should be run periodically as cron jobs.

- Retrieval of forwarding tables and ARP caches for IP/MAC address tracking
- Devices should be re-discovered via SNMP frequently to maintain an accurate list of ports, ip addresses, etc.
- Rediscovery of network topology
- Netdot keeps history records for some objects every time they are updated. With time, old history should be deleted from the database to save disk space.
- Netdot can generate text documentation that is easy to find via simple grepping commands, for example, information about people, locations, device port assignments, etc. This documentation should be kept up to date by exporting it frequently.
- Configurations for external programs can be generated using Netdot data. Current supported programs include: Nagios, Sysmon and Rancid.

- The `netdot.cron` file included in the package is a sample crontab containing recommended periodic jobs. You should customize it to your liking and copy it to your cron directory, for example:

```
# cp netdot.cron /etc/cron.d/netdot
```

Operation

Device Management

Netdot can discover and maintain an extensive amount of information about network devices. The easiest way to gather and store this information is by querying the devices using the Simple Network Management Protocol (SNMP). Devices can be discovered individually, by subnet, or by providing a text file with a list of devices.

Device Discovery using the web UI

Go to Management -> Devices. In the Tasks section, type the hostname or IP address of the device in question, along with the SNMP community and click [discover]. Netdot will then query the device using SNMP and present a window where you can assign an owner entity (for example, your organization), the entity that uses the device (for example, your customer), the location and a contact list.

If you are discovering a layer 3 device with ip forwarding turned on (such as a router or firewall), Netdot will ask you if you would like to automatically create subnets, based on the IP configuration of the device interfaces. This is a convenient way to add your subnets to Netdot.

Another option is to specify whether Netdot should assign any newly created subnets the same owner and user entities assigned to the device.

Once you click the [update] button, Netdot will show the discovery information and a link to the device page at the button.

You can always re-discover a device manually by using the [snmp-update] button on the top right corner of the device page. For example, if you have added a new port adapter, new interface cards, or if the device has been replaced with a different unit.

Device discovery using the command line interface (CLI)

For brevity, let's assume you are located at the Netdot installation prefix, for example, `/usr/local/netdot`.

You can discover a single device by executing:

```
bin/updatedevices.pl -H <device-name> -I -c <community>
```

You can also try discovering a whole subnet like this:

```
bin/updatedevices.pl -B 192.168.1.0/24 -I -c <community>
```

In addition, you can give Netdot a specific list of devices you would like to discover:

```
bin/updatedevices.pl -F <text-file> -I -c <community>
```

The file should contain a list of devices and their communities, separated by spaces, one device/community per line, for example:

```
device1 community1
device2 community1
device3 community2
...
```

Netdot can retrieve ARP caches and bridge forwarding tables. You will probably want to fetch ARP caches from your layer 3 devices (i.e. routers and firewalls), and forwarding tables from your layer 2 devices (switches). Examples:

```
bin/updatedevices.pl -H <router> -A -c <community>
```

```
bin/updatedevices.pl -H <switch> -F -c <community>
```

Netdot can also try to discover the network topology. For that, you need to run:

```
bin/updatedevices.pl -T
```

Ideally, once you have discovered all your devices, you should combine all this functionality and have it run periodically via CRON. A sample crontab entry would be:

```
0 * * * * root /usr/local/netdot/bin/updatedevices.pl -I -A -F -T
```

You will find some examples of cron jobs in the file named *netdot.cron*

Device Documentation

Once you have created a device, you can go ahead and add more information about it.

Going to Management -> Devices you can search a device based on its name, IP or MAC address.

From the device page, you can navigate to the different sub-sections depending on the information you want to edit.

Interfaces: Here you can edit interface descriptions, assign network jacks, etc. by clicking on the [edit] button. You can also edit a specific interface by clicking on its number or on its name. If you are running topology discovery, you will probably see neighbor information. If for some reason the topology discovery process cannot detect a neighbor, you can add it manually by clicking on the [add] button in the neighbor column.

Manually adding a neighbor sets the "Neighbor Fixed" flag in the Interface object. This flag prevents the topology discovery process from removing the neighbor relationship.

Tip

Neighbor relationships tend to change frequently as hardware is replaced and connections are moved. Therefore, fixed neighbor settings can become out of date pretty soon. It is preferable to let the topology discovery process maintain neighbor relationships.

IP Address Space

Netdot is a useful tool for IP address space management (referred to as *IPAM* these days). Some key features are:

- Address space is hierarchically organized through the use of a fast binary tree algorithm, which is the same technique used by routers for doing route lookups.

- New subnets can be automatically created based on the interface information retrieved from routers and firewalls.
- Easy visualization of used vs. available address space for easier allocations

IP blocks

IP objects are called IP blocks. These objects can represent individual end-node addresses, as well as groups of addresses. The distinguishing characteristic is the prefix attribute. For example, an IPv4 block with a 32 bit prefix is an end-node address, while a block with a 24 prefix represents a group of 254 end-node addresses.

Each address or block has a corresponding status. Let's see those in detail.

IP block Status

IP objects are assigned a status to better represent their nature. Depending on whether the IP is an end address or a block, different status values can be assigned.

The status of an end-node address can be one of:

- Static: These are addresses that have been statically assigned to hosts or device interfaces.
- Dynamic: Addresses that belong to a DHCP pool
- Discovered: Addresses that have not been assigned as static or dynamic, but have been seen on the network (as part of ARP entries, for example).
- Reserved: Addresses that should not be assigned

On the other hand, the the status of an IP block can be one of:

- Container: This kind of block is meant to group or contain other blocks, such as Subnet blocks or other Container blocks. For example, let's say your whole IPv4 address space is 192.168.0.0/16. You also have partitioned this space into two /17 blocks, and from these blocks, you allocate subnets that you configure in your routers. In this case, you would have:

```
192.168.0.0/16 -> Container
  192.168.0.0/17  -> Container
    192.168.0.1/24 -> Subnet
    192.168.0.2/24 -> Subnet
  192.168.128.0/17 -> Container
    192.168.128.10/24 -> Subnet
    192.168.128.20/24 -> Subnet
```

- Subnet: This kind of block is meant to represent actual subnets that are configured on the interfaces of your layer 3 devices such as routers or firewalls. Subnets usually contain the end-node addresses that you assign to your users.
- Reserved: Similarly to reserved addresses, reserved blocks are not supposed to be allocated for whatever reason.

DNS

Netdot can maintain DNS zone data. Zones can be exported as text files to be used by DNS server software. Currently, only ISC BIND zone file exporting is supported.

Tip

The mechanisms by which zone files are transferred to and loaded by authoritative name servers are left to the administrator. A simple way to do this is by running a name server locally in the machine that runs netdot, and saving those zone files in the location where the software can load them periodically. A more complex setup could involve saving these files into revision control systems (CVS, SVN, etc), which could then be used by system configuration engines like Puppet or CfEngine to run syntax checks and load them into the appropriate name servers.

Netdot supports the following DNS records: A, AAAA, CNAME, DS, HINFO, LOC, MX, NAPTR, SRV, and TXT.

You can import your existing BIND zones into netdot with the help of the tool *import_bind_zones.pl* from the *import* subdirectory.

To add a new zone manually, go to Management -> Zones and click on [add]. You will see a new zone created with default values extracted from the configuration file.

Once a zone is created, it must be linked to an IP block (Subnet or Container). You can do this by clicking on the [add] button of the IP blocks section in the zone page.

At this point, you can add new records by clicking on the [add] button on the Records section. Records can also be added from other parts of the user interface, for example, from the IP address page, or the Hosts page.

Records can also be imported in bulk into the zone by clicking on the [import] button of the Records section and pasting the text from a BIND zone file into the text box.

Each time the zone or its contents are modified, the transaction is added to a list of "pending changes". This list is used to determine when a zone needs to be exported. Zones can be exported manually via the UI by going to the Export menu, or via cron jobs. When a zone is exported, its serial number is increased and its list of pending changes is flushed.

Netdot can auto-generate and maintain DNS records for IP address belonging to device interfaces, for example, routers. The logic of how these names are generated is handed off to a plugin.

The '@' record

In netdot, as in BIND, the '@' record symbolizes the domain. In order to add records that apply to the domain itself, such as NS records, you must add the '@' name first. Once this name is created, you can add NS, MX and A records to it, which will apply to the domain that the zone represents.

DHCP

Netdot can maintain DHCP information and generate configurations for ISC DHCPD.

DHCP information is organized hierarchically around the DHCP Scope object. Netdot supports scope definitions of the following types: global, subnet, shared-subnet, group, and host. Each of these scopes can be assigned one or more attributes.

Contact Information

Cable Plant

Advanced DB operations

Reports

Exporting Configurations for External Programs

You can use the exporter tool to generate text files that can be used as configurations for third-party tools and programs.

The exporter tool is available in the web UI, under the "Export" tab. Simply select one or more programs and click on the [submit] button. Netdot will show some output from the exporter tool, including the paths to the new files.

Additionally, the exporter can be called from the command line. For example, to generate Nagios configurations:

```
bin/exporter.pl -t Nagios
```

Or you can export several in one call:

```
bin/exporter.pl -t Nagios,Sysmon,Rancid,BIND,DHCPD
```

There are specific export parameters for each of these which you can customize by editing your Site.conf file.

Authorization

Starting with version 0.9, netdot supports multi-level user access.

There are three types of users that correspond with levels of access in Netdot:

- Admin: Full access to the UI and operations on objects.
- Operator: Full access to the UI, but read-only access to objects.
- User: Limited UI, with view, edit, and delete access to particular objects.

Assigning permissions to unprivileged users

Permissions can be assigned to individuals or to groups. Individuals are grouped in contact lists. A user who is a member of a contact list inherits the permissions from the list. However, the individual can have more specific permissions (or no permissions) if necessary.

There is a limited number of objects which unprivileged users can gain access to:

- DNS records: Users can create, modify and delete records from a certain zone. Permissions can be given for the entire zone or for subsets of it, based on IP blocks. For example, if a user is given view, edit and delete permissions to myzone.com, he or she can view, modify and remove any record from that zone. On the other hand, if the zone covers hosts from a supernet, i.e. 10.0.0.0/16, and the user should only have control on records within a particular subnet, i.e. 10.0.0.0/24, instead of assigning permissions on myzone.com, the administrator can assign view, edit and delete permissions on that particular subnet.
- Device interfaces: Users can view port details such as number, name, vlan, room, jack, description and neighbor. A user can only edit the room, jack and description fields. To assign permissions to a

user on a list of devices, select the Device class and then select one or more devices to which the user can have access.

- Contact Lists: A user can add, modify and delete contacts from given contact lists.

To assign permissions for an individual user, perform the following tasks:

- Make sure there is a Person object for the user. You can verify if a Person object exists by going to Contacts -> People and searching for the person's name in the Search box. If the object does not exist, you can create a new one by clicking on the [new] button on the upper right corner of the same window.
- Make sure that the person object has a Username and a User Type set. If you have configured netdot to use Radius or LDAP authentication, make sure that the username corresponds with the login information in those central authentication systems. If you are using local authentication instead, make sure that you set a local password using the Password field.
- On the Person page, you can add permissions by clicking on the [access_rights] button. This will display current permissions. You can now add new ones by clicking on the [add] button on the right.
- On the UserRight window, select the Object Class, the specific object or objects, and one or more access rights (view, edit, delete). Only select the 'none' right to revoke all permissions inherited from a group. Click on 'Insert'.